

# Explaining single prediction

Mateusz Staniak

# Why?

- ▶ to increase *trust* in important predictions
- ▶ to *understand* which factors drive the prediction
- ▶ to inspect unusual/outlier predictions
- ▶ to improve the model

How?




Figure 1:

Now details!

LIME

## LIME: general idea




Figure 2:

## LIME: how it's done

- ▶ Gaussian sampling for tabular, uniform sampling from interpretable inputs for image/text.
- ▶ Scores for new observation are weighted by the distance from original observation.
- ▶ Variable selection is usually based on ridge/lasso regression.
- ▶ Weights are assigned to interpretable inputs to decide if they *vote* for or against a given label.

## LIME: summary

LIME approximates the complex model with a simple model and discovers which features in *interpretable representation* of data drive the prediction.

Pros:

- ▶ intuitive & suitable for different types of data
- ▶ easy to understand & interpret

Cons:

- ▶ results can be inconsistent (see: Lundberg)
- ▶ depends on many *hyperparameters* (kernel, size, ...)



## Shapley Values

## Shapley values: general idea

- ▶ The goal is a decomposition of prediction into sum of scores related to (simplified) features.
- ▶ The problem is solved using game theory: *Shapley values*.
- ▶ This approach unifies several methods (including LIME).

## Shapley values: how it's done

- ▶ Different approximate algorithms were proposed.
- ▶ Exact methods exist for linear models and tree ensemble methods.
- ▶ Classic way: sample permutations of variables, then average contributions.
- ▶ Better way: approximation based on LIME and Shapley values for regression.

## Shapley values: summary

Shapley values decomposition prediction into feature contributions in a rigorous way rooted in game theory.

Pros:

- ▶ good theoretical properties
- ▶ comes with good visual diagnostic tools

Cons:

- ▶ won't produce sparse explanations
- ▶ assumes *additivity*, uses simplified inputs (exception: trees)
- ▶ computational issues

LIVE

## LIVE: general idea

- ▶ Modification of LIME for tabular data and regression problems.
- ▶ Similar observation for *fake* dataset are sampled from empirical distributions.
- ▶ Variable selection is possible.
- ▶ Focused on model visualization.

## LIVE: how it's done

- ▶ All new observations are treated as similar (identity kernel & equal weights) (Because of the use of empirical distributions)
- ▶ Variables are selected by LASSO (`glmnet` package).
- ▶ All models from `mlr` are supported both as black box and white box.
- ▶ Model visualization for linear models in particular.

## LIVE: summary

live fits a simple model to approximate the black box locally and allows model visualization for this model.

Pros:

- ▶ flexible & focused on model visualization
- ▶ local exploration in data space (no binary inputs)

Cons:

- ▶ comes with no theoretical guarantees (though *works* at least in simple cases and is pretty stable)
- ▶ inherits strengths/weaknesses of white box model



breakDown

## breakDown: general idea

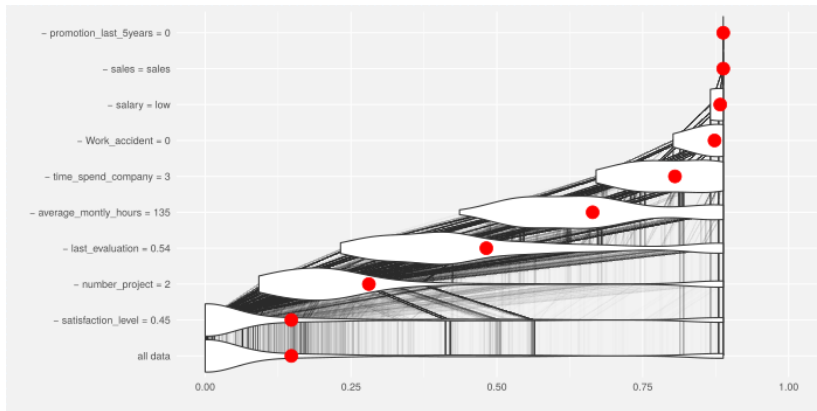


Figure 3:

## breakDown: how it's done

- ▶ Contributions are assigned in a greedy way.
- ▶ Waterfall plots are important tool as a visual representation.

breakDown: summary

breakDown decomposes the prediction into feature contributions and visualizes them.

Pros:

- ▶ simple computation
- ▶ easy to interpret (in particular compared to SHAP)

Cons:

- ▶ limited to additive effects
- ▶ again, no theoretical guarantees (though can be thought of as a rough Shapley values estimate)

Let's see some examples!

# Dataset

```
library(DALEX)
library(live)
library(lime)
library(shapleyr)
library(mlr)
library(tidyverse)
library(e1071)
load("wine.rda")
glimpse(wine)
```

```
## Observations: 1,599
```

```
## Variables: 12
```

```
## $ fixed_acidity      <dbl> 7.4, 7.8, 7.8, 11.2, 7.4, 7.4, 7.8, 11.2, 7.4, 7.4, 7.8, 11.2
```

```
## $ volatile_acidity  <dbl> 0.700, 0.880, 0.760, 0.280, 0.260, 0.290, 0.390, 0.390, 0.390, 0.390, 0.390, 0.390
```

```
## $ citric_acid        <dbl> 0.00, 0.00, 0.04, 0.56, 0.00, 0.00, 0.00, 0.00, 0.00, 0.00, 0.00, 0.00
```

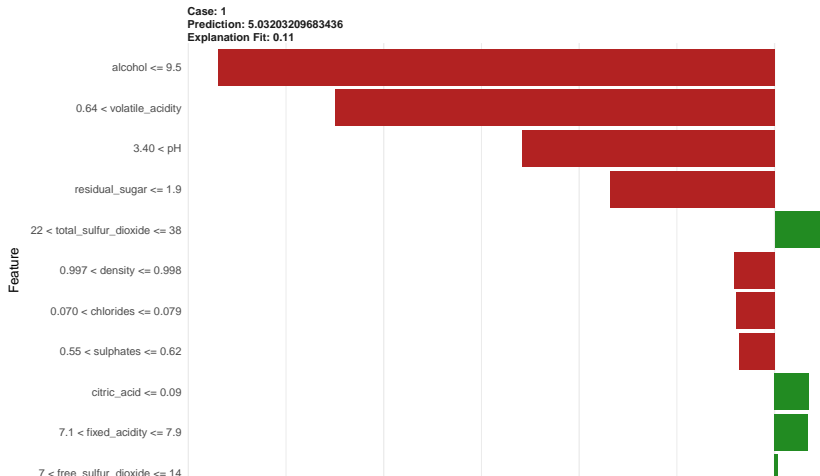
```
## $ residual_sugar    <dbl> 1.9, 2.6, 2.3, 1.9, 1.9, 1.9, 1.9, 1.9, 1.9, 1.9, 1.9, 1.9
```

```
## $ chlorides          <dbl> 0.076, 0.098, 0.092, 0.075, 0.076, 0.076, 0.076, 0.076, 0.076, 0.076, 0.076, 0.076
```

```
## $ free_sulfur_dioxide <dbl> 11. 25. 15. 17. 11. 13. 15. 11. 11. 11. 11. 11.
```

# LIME

```
lime_obj <- lime(wine, wine_svm)
model_type.svm <- function(x, ...) "regression"
lime_expl <- lime::explain(wine[nobs, ], lime_obj, n_features = 13)
plot_features(lime_expl)
```



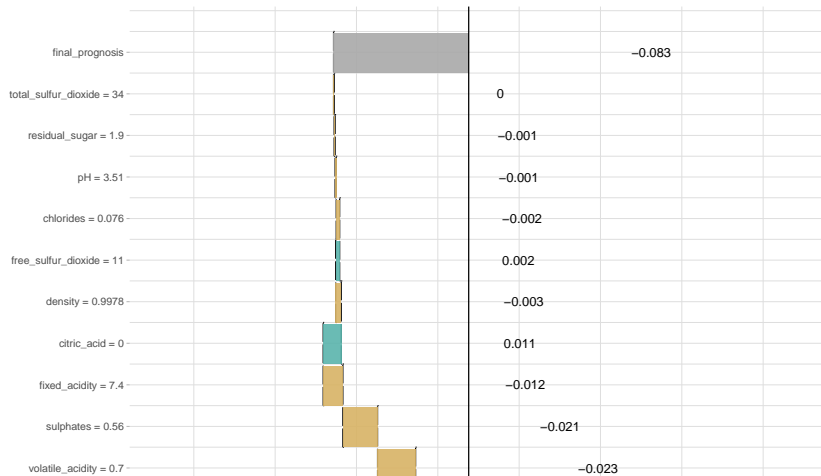
## Shapley values

```
tsk <- makeRegrTask(data = wine, target = "quality")
model <- train("regr.svm", tsk)
shap_expl <- shapley(nobs, model = model, task = tsk)
gather(shap_expl, "variable", "contribution") %>%
  arrange(desc(abs(contribution)))
```

##	variable	contribution
## 1	volatile_acidity	-0.232
## 2	alcohol	-0.207
## 3	sulphates	-0.199
## 4	citric_acid	0.169
## 5	fixed_acidity	-0.085
## 6	pH	-0.070
## 7	free_sulfur_dioxide	0.034
## 8	residual_sugar	-0.020
## 9	total_sulfur_dioxide	0.015
## 10	chlorides	0.012
## 11	density	0.001

# LIVE

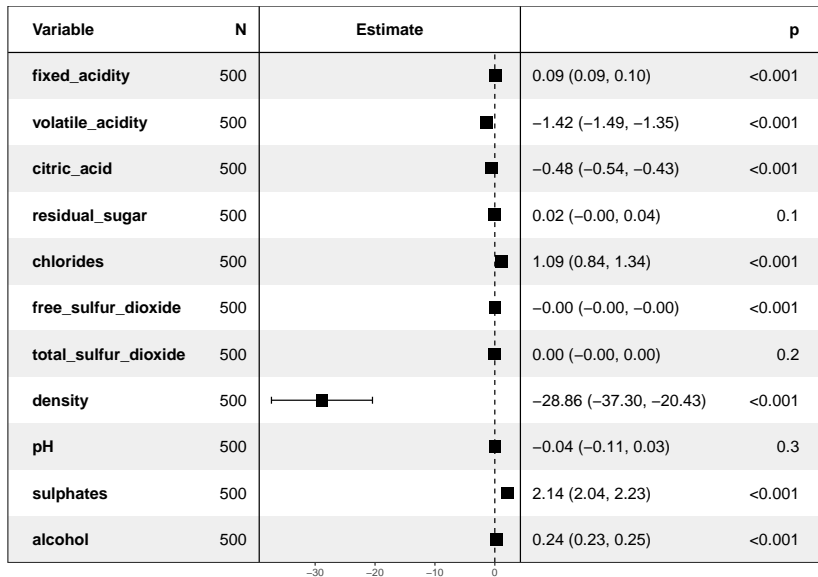
```
local <- sample_locally(wine, wine[nobs, ], "quality", 500)
local <- live::add_predictions(wine, local, wine_svm)
live_expl <- fit_explanation(local, "regr.lm")
plot_explanation(live_expl, "waterfallplot", wine[nobs, ])
```






## LIVE cont.

```
plot_explanation(live_expl, "forestplot")
```



## breakDown

```
wine_explainer <- DALEX::explain(wine_svm, data = wine, y =  
wine_expl <- single_prediction(wine_explainer, wine[nobs, ]  
plot(wine_expl)
```



Time for discussion!