

Prediction understanding

Mateusz Staniak

Outline

1. Intro & **breakDown** package.
2. Local approximations: **live** package.
3. Summary & state of the art solutions (LIME and Shapley values)

Intro & prediction breakdown

Why explain a single prediction?

(Bird's-eye view)

- when important decision are made based on ML model, it needs to be **trustworthy**
- trust comes from **understanding**
- the demand for interpretable algorithms is growing (see: *Weapons of math destruction*, Facebook feed controversies etc.)

(Worm-eye view)

- this demand is transfered into legal regulations (see: RODO)

=> more and more institutions have to explain model predictions (debt collection, loans ...)

- understanding models helps improve them

Which predictions need explanation?

1. Every prediction the client (or the boss) wants to understand
2. Predictions that seem suspicious
 - How to spot them?
 - How to explain them?




Figure 1:



Figure 2:

=> model performance

=> model diagnostics


Model diagnostics: example data

```
library(tidyverse)
library(randomForest)
load("./rda_files/houses.rda")
houses

## # A tibble: 5,851 x 7
##   rooms_num area sqm_price year floor max_floor district
##   <int> <dbl> <int> <int> <int> <int> <fct>
## 1      3  89.0   5270  2007     2      2 Krzyki
## 2      4 163.   6687  2002     2      2 Psie Pole
## 3      3  52.0   6731  1930     1      2 Srod miescie
## 4      4  95.0   5525  2016     1      2 Krzyki
## 5      4  88.0   5216  1930     3      4 Srod miescie
## 6      2  50.0   5600  1915     3      4 Krzyki
## 7      2  48.0   9146  2010     2      6 Fabryczna
## 8      2  46.0   9761  2011     4      6 Stare Miasto
## 9      3  67.0   5209  1990     1     10 Fabryczna
## 10     2  45.8   6878  2000     3      4 Krzyki
## # ... with 5,841 more rows

hrf <- randomForest(sqm_price ~., data = houses)
```


Model diagnostics: predicted vs observed values



- Points on the plot should be close to the $y = x$ line,
- Questions:
 - is there a pattern? (For example: does the deviation from true value grow as true value grows?)
 - are there any points especially far from the line (meaning: points with large residuals)?
- More diagnostic tools: `auditor` package

Model performance

```
library(DALEX)
rf_explainer <- explain(hrf, data = houses, y = houses$sqm_price)
rf_perf <- model_performance(rf_explainer)
plot(rf_perf, geom = "boxplot")
```



- shortly summarizes the distribution of the absolute value of residuals
- red dot is the root mean square error
- we can put boxplots for several models on the same plot (simply by passing the as arguments to `plot`) to compare models
- boxplots help discover outliers

Single prediction explanation

- Once we identified predictions we want to explain, we need tools that will help us!
- Methods:
 - LIME
 - Shapley values
 - Break Down
 - LIVE
- Two main ideas:
 - Attribute scores to explanatory variables according to their influence on the prediction (**contributions**)
 - Fit a model locally around an observation and investigate it
 - NOTE: both approaches lead to local feature importance

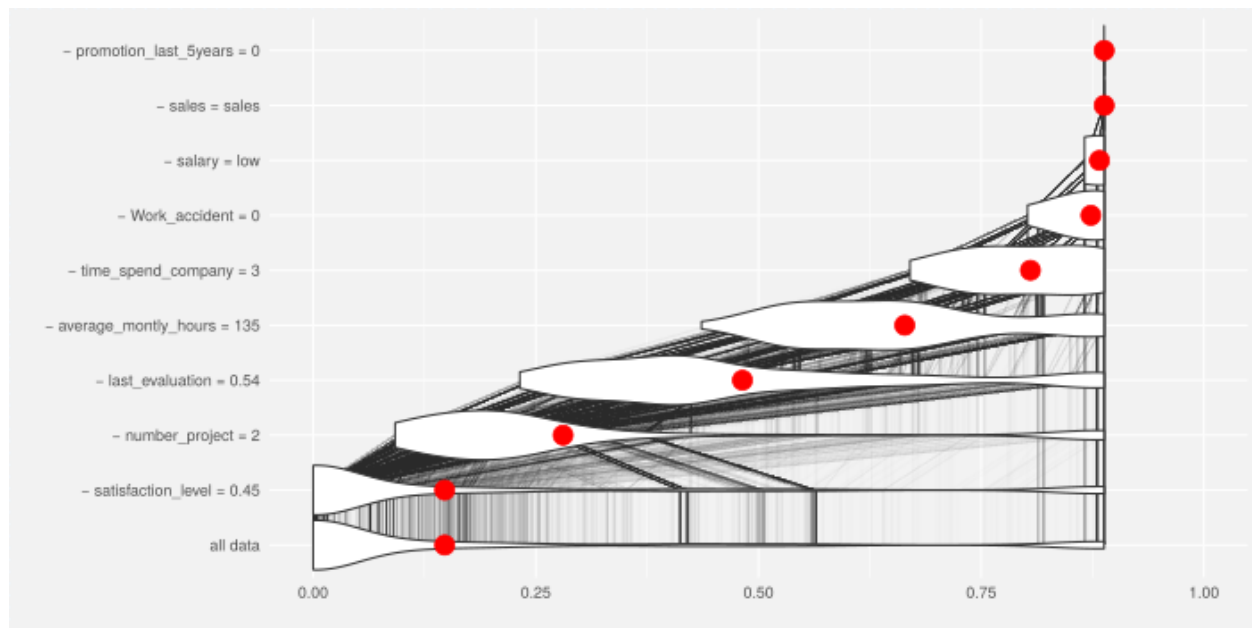


Figure 3:

- Contributions: Shapley values and Break Down
- Local models: LIME and LIVE

Break Down

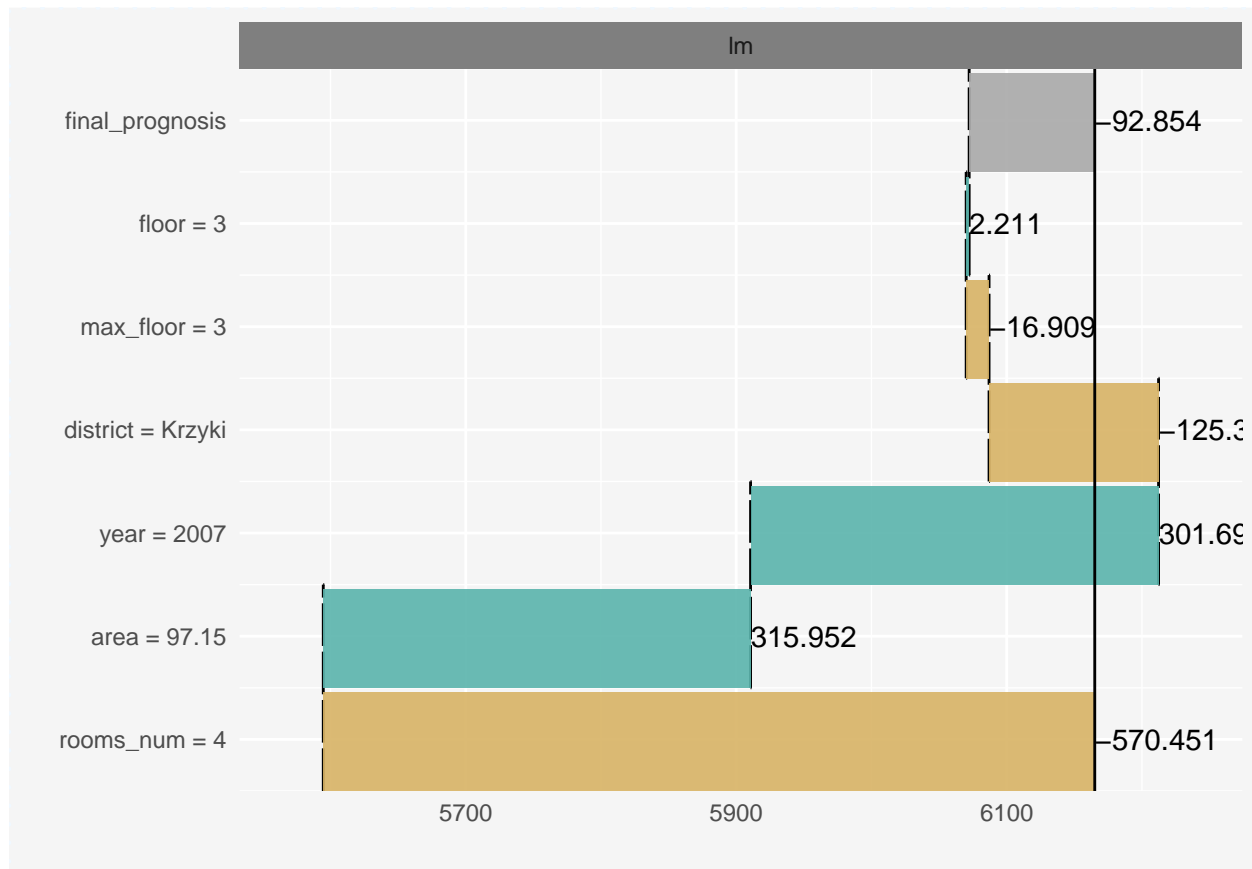
- General idea
- Another approach to finding *additive* feature contributions
- Contributions are assigned in a greedy manner
- Waterfall plots as a visual tool

=> **more intuitive interpretation**

Example


- Break Down for linear models

```
linear_model <- lm(sqm_price ~., data = houses)
lm_explainer <- DALEX::explain(linear_model, data = houses, y = houses$sqm_price)
breakdown_linear <- single_prediction(lm_explainer, houses[4036, -3])
plot(breakdown_linear)
```



- Contributions are scaled, so they do not depend on the scale of the data (insensitive to location/scale)
- We can see actual contributions, not just the weights (as in LIME plots)
 - Model-agnostic Break Down


```
breakdown_explanation <- single_prediction(rf_explainer, houses[4036, -3])
plot(breakdown_explanation)
```



- We can see how important District and Year are in this random forest prediction

But isn't it enough to calculate feature importance?

```
global_feat_imp <- DALEX::variable_importance(rf_explainer)
plot(global_feat_imp)
```

- No. Particular instances can be influenced the most by different features, not necessarily the ones that are most important globally.

Time to practice!

First set of exercises goes here

Part II: local approximations

LIVE (Local Interpretable Visual Explanations)

- General idea
 - Modification of **LIME** for tabular data and regression problems with emphasis on model visualization.
 - Similar observations for *fake* dataset are sampled from empirical distributions.
 - Variable selection is possible (LASSO, then explanation model is fitted to selected features).
- More details
 - Two methods of creating the new dataset are available: by permuting each variable and by changing one feature per observations
 - We can control which variables are allowed to vary through **fixed** variable argument to **sample_locally** (keeping date/factor/correlated variables unchanged)

Example

```
library(live)
library(mlr)
new_dataset <- sample_locally2(data = houses,
                               explained_instance = houses[4036, ],
                               explained_var = "sqm_price",
                               size = 1000)
with_predictions <- add_predictions2(new_dataset, hrf)
live_explanation <- fit_explanation2(with_predictions, "regr.lm")
live_explanation
```

```
## Dataset:
## Observations: 1000
## Variables: 7
## Response variable: sqm_price
## Explanation model:
## Name: regr.lm
## Variable selection wasn't performed
## Weights present in the explanation model
## R-squared: 0.9889
```

Aktualnie standaryzacja zmiennych jest niepotrzebnie w `sample_locally`, jak to zmienię, `forestplot` będzie lepiej wyglądał, bo będzie po standaryzacji

- Default method of sampling is *live*, default explanation model is linear regression and distance is measured (weights are assigned) by gaussian kernel.
- Plot local model structure: forest plot

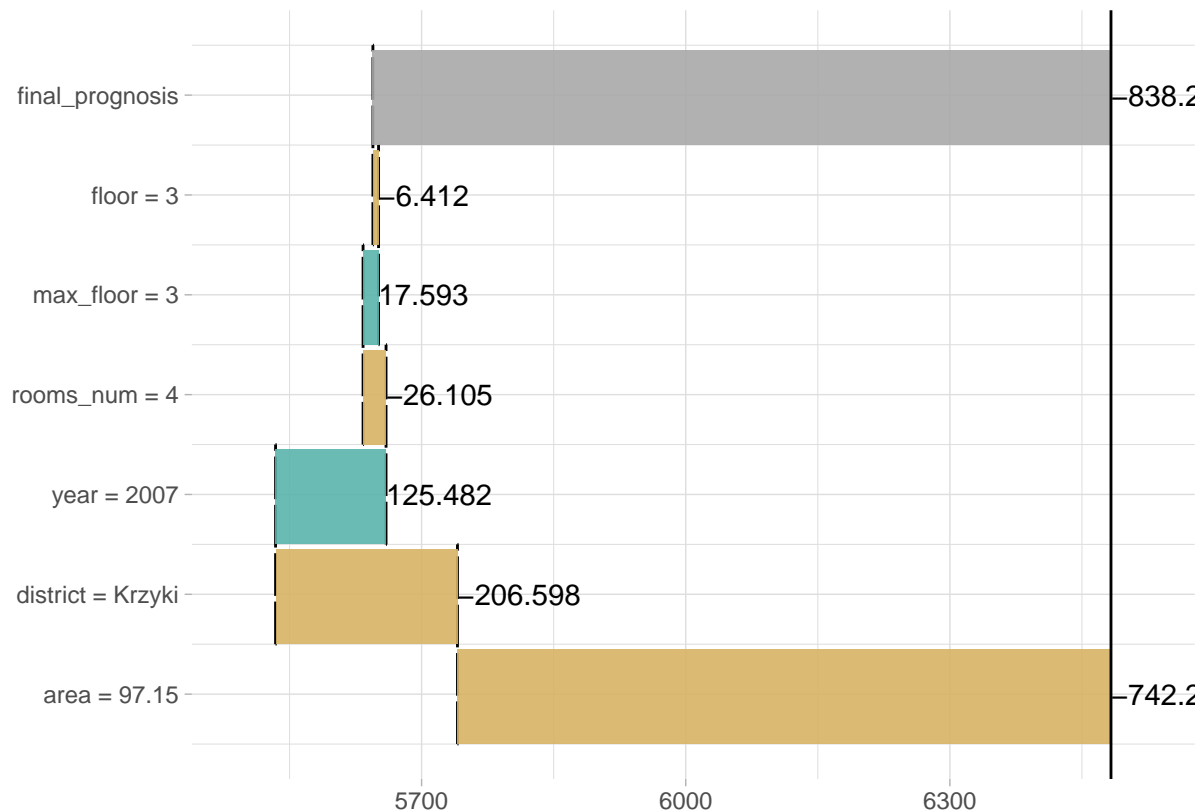
```
plot_explanation2(live_explanation, "forest")
```

Variable	N	Estimate	p
rooms_num	1000	■ -118.12 (-135.56, -100.68)	<0.001
area	1000	■ -144.53 (-270.30, -18.75)	0.024
year	1000	■ 39.25 (-8.76, 87.26)	0.109
floor	1000	■ -80.15 (-99.78, -60.52)	<0.001
max_floor	1000	■ -150.37 (-170.09, -130.65)	<0.001
district	Fabryczna 41	■ Reference	
	Krzyki 886	■ 39.67 (14.55, 64.80)	0.002
	Psie Pole 19	■ -571.74 (-635.10, -508.38)	<0.001
	Srodmiescie 30	■ 1749.10 (1696.60, 1801.61)	<0.001
	Stare Miasto 24	■ 7062.97 (7005.53, 7120.41)	<0.001
(Intercept)		■ -57964.87 (-155073.63, 39103.90)	

-15000000000

- Plot local variable contributions: waterfall plot (Break Down)

```
plot_explanation2(live_explanation, "waterfall")
```



Time to practice!

Second set of exercises goes here

Part III: state-of-the-art solutions & summary

LIME (Locally Interpretable Model-agnostic Explanations)

- General idea
- Some details:
 - Gaussian sampling for tabular, uniform sampling from interpretable inputs for image/text.
- Scores for new observation are weighted by the distance from original observation.
- Variable selection is usually based on ridge/lasso regression.
- Weights are assigned to interpretable inputs to decide if they *vote* for or against a given label.
- Note: method depends on many hyperparameters

Example

```
library(mlr)
house_task <- makeRegrTask(data = houses, target = "sqm_price")
house_rf_mlr <- train("regr.randomForest", house_task)
```




Figure 4:

```
library(lime)
explained_prediction <- houses[4036, ]
lime_explainer <- lime(houses,
                      model = house_rf_mlr)
lime_explanation <- lime::explain(houses[4036, ],
                               explainer = lime_explainer,
                               n_features = 5)
```

- weights from ridge regression are on the plot (NOT weights multiplied by actual feature values)
- positive weights are *for*, negative weights are *against*

Shapley values


- General idea
 - The goal is a decomposition of prediction into a **sum** of scores related to (simplified) features.
 - The problem is solved using game theory: *Shapley values*. Variables are *players* who contribute to the outcome - the prediction - and we try to *pay* them accordingly to their contributions.
- This approach unifies several methods (including LIME).
- Some details
 - Exact methods exist for linear models and tree ensemble methods. In other cases, approximations are needed.
 - The classic way: sample permutations of variables, then average contributions.
 - The better way: approximation based on LIME and Shapley values for regression.
- This method has good theoretical properties, but will not produce sparse explanations

Example

```
library(shapleyr)
shapley_explanation <- shapley(4036,
                             task = house_task,
                             model = house_rf_mlr)
class(shapley_explanation) <- c("shapley.singleValue", "list")
gather(shapley_explanation$values, "feature", "shapley.score")
```

```
##      feature shapley.score
## 1      _Id      4036
## 2     _Class      <NA>
## 3 rooms_num    -93.112
## 4      area   -311.833
## 5      year    527.284
## 6      floor  -139.601
## 7 max_floor    53.036
## 8 district   -622.128
```

```
plot(shapley_explanation)
```



- **0** is the mean of all predictions
- the **black dot** is the prediction we are explaining
- values and the plot describe how we move from the global mean of predictions to this particular predictions
- most important features are the ones that help move the most

Acknowledgement

Podziękowanie dla UW...

Time to practice!

Third set of exercises goes here