

eRum exercises

Mateusz Staniak

Exercises

Exercise 1

Run the following code to fit random forest, linear regression and SVM to the housing prices data.

```
library(tidyverse)
library(live)
library(DALEX)
library(randomForest)
library(e1071)
library(auditor)
load("./rda_files/houses.rda")

set.seed(33)
house_rf <- randomForest(sqm_price ~., data = houses)
house_svm <- svm(sqm_price ~., data = houses)
house_lm <- lm(sqm_price ~., data = houses)
```

Create DALEX explainer object for each of the models. Create and compare boxplots of residuals for all the models (`model_performance`).

- Which model is the best?
- Are there any outlying predictions?
- Find the observation with the largest absolute value of residual among houses cheaper than 7000 PLN.

TIP: object returned by `model_performance` function is a data frame with colnames *predicted*, *observed*, *diff* and *label*.

Exercise 2

Create single prediction explainers for the instance chosen in **Exercise 1**. Create Break Down plots for each of the observations. What are the keys factors that drive the prediction? Are they the same for every model?

Exercise 3

Run the following code to train model random forest model using `mlr` interface (this is necessary for `shapleyR` package).

```
library(mlr)
load("./rda_files/houses.rda")
n_obs <- 1189

house_task <- makeRegrTask(data = houses, target = "sqm_price")
house_rf_mlrl <- train("regr.randomForest", house_task)
```

Use `shapleyR` package to calculate Shapley values for prediction chosen in **Exercise 1** (its index is in `n_obs` object). Are the results consistent with Break Down results from **Exercise 2**? Draw a plot of Shapley values.

TIP: remember to set class of the object returned by `shapley` function to `shapley.singleValue` before using `plot`.

Bonus 1:

Draw plots of fitted vs observed values for each of these models. Can you spot any problems with the predictions? Are the prices usually underestimated or overestimated?

Bonus 2:

Create variable importance explainer. Compare global variable importance to scores obtained in **Exercise 2** and **Exercise 3**.

Exercise 4

```
n_obs <- 1189
```

Simulate new data around the observation from **Exercise 1** (its index is in the `n_obs` object.) and the add random forest predictions. Then fit a linear model locally.

TIP: remember to load `mlr` package. TIP2: don't use too small `size` for the simulated dataset. I recommend at least 1000.

Exercise 5

Visualize approximation created in **Exercise 4**. Use `plot_explanation2` function to create a forest plot of the linear model and then the Break Down plot.

Exercise 6

Use `lime` package to approximate random forest model around prediction chosen in **Exercise 1** (its index is in `n_obs` object). Follow the `lime` work flow: 1. Create an explainer. 2. Approximate the model around the explained instance. 3. Use `plot_features` function to see, how features influence this prediction.

TIP: use `house_rf_mlr` object from **Exercise 5**, because `lime` works well with `mlr` objects.

Bonus 3

Use `add_predictions` function to add SVM and LM predictions to the simulated dataset. Compare plots for all three models.

Bonus 4

Run the following code to see largest residuals for *Psie Pole* district.

```
library(tidyverse)
houses %>%
  mutate(id = 1:n()) %>%
  mutate(rf_pred = predict(house_rf)) %>%
  mutate(abs_res = abs(sqm_price - rf_pred)) %>%
```

```

arrange(desc(abs_res)) %>%
filter(sqm_price < 7000,
       district == "Psie Pole") %>%
head(5)

```

```

## # A tibble: 5 x 10
##   rooms_num area sqm_price year floor max_floor district    id rf_pred
##   <int> <dbl>   <int> <int> <int>   <int> <fct>   <int> <dbl>
## 1      2  46.0    2174  2005     2       5 Psie Pole  5830  6541.
## 2      2  46.0    1957  2001     4       5 Psie Pole   942  5829.
## 3      3  65.0    1877  2001     8      10 Psie Pole  4303  5663.
## 4      3  75.0    4267  2013     0       0 Psie Pole  4308  6388.
## 5      4  88.4    3394  2004     4       5 Psie Pole  3489  5495.
## # ... with 1 more variable: abs_res <dbl>
n_obs2 <- 5830

```

Using `live` package, fit a linear model around the top observation. Compare waterfall plots for this prediction and the prediction from **Exercise 5**. How are they different?

Solutions

Exercise 1

```

rf_expl <- DALEX::explain(house_rf, data = houses, y = houses$sqm_price)
svm_expl <- DALEX::explain(house_svm, data = houses, y = houses$sqm_price)
lm_expl <- DALEX::explain(house_lm, data = houses, y = houses$sqm_price)

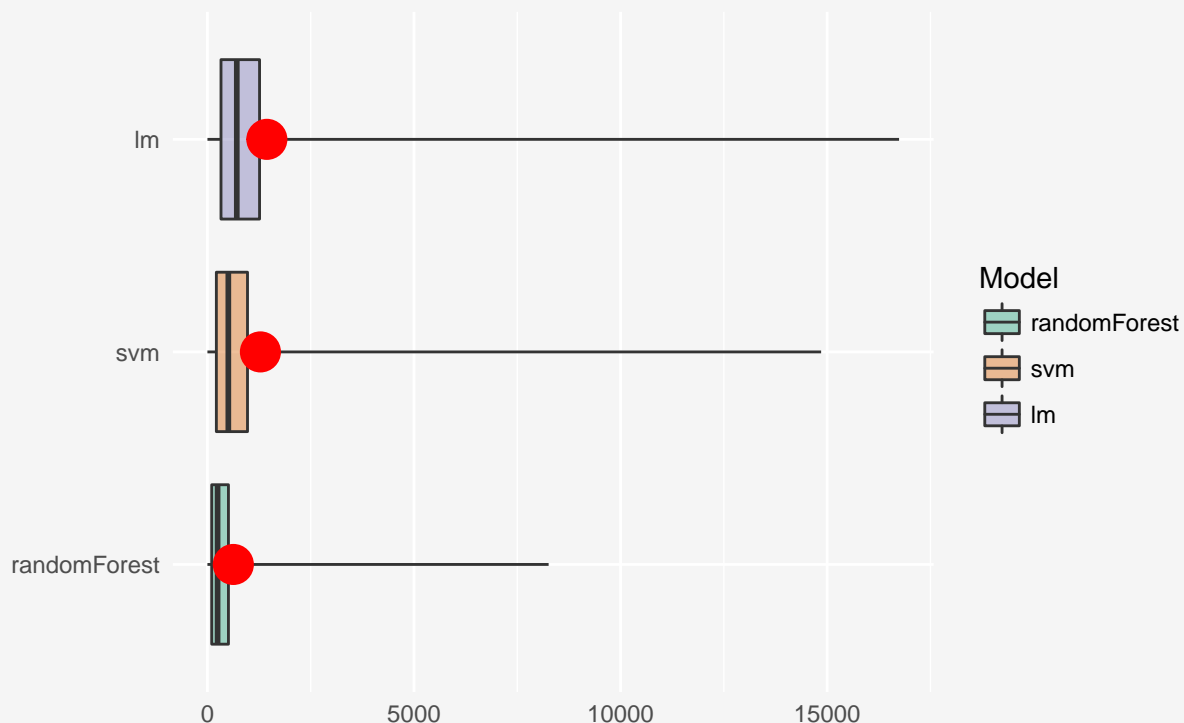
rf_perf <- model_performance(rf_expl)
svm_perf <- model_performance(svm_expl)
lm_perf <- model_performance(lm_expl)

plot(rf_perf,
     svm_perf,
     lm_perf,
     geom = "boxplot")

```

Boxplots of | residuals |

Red dot stands for root mean square of residuals



```
rf_perf %>%
  mutate(id = 1:n()) %>%
  arrange(desc(abs(diff))) %>%
  filter(observed < 7000) %>%
  head(5)
```

##	predicted	observed	diff	label	id
## 1	4338.774	1585	2753.774	randomForest	1189
## 2	4842.923	2174	2668.923	randomForest	5830
## 3	7399.405	4750	2649.405	randomForest	4824
## 4	6326.248	3742	2584.248	randomForest	4419
## 5	8345.797	6000	2345.797	randomForest	4005

```
svm_perf %>%
  mutate(id = 1:n()) %>%
  arrange(desc(abs(diff))) %>%
  filter(observed < 7000) %>%
  head(5)
```

##	predicted	observed	diff	label	id
## 1	5771.421	1585	4186.421	svm	1189
## 2	7983.699	4063	3920.699	svm	356
## 3	5722.197	1957	3765.197	svm	942
## 4	5907.319	2174	3733.319	svm	5830
## 5	5544.827	1877	3667.827	svm	4303

```
lm_perf %>%
  mutate(id = 1:n()) %>%
```

```

arrange(desc(abs(diff))) %>%
filter(observed < 7000) %>%
head(5)

```

```

##   predicted observed    diff label   id
## 1  6272.608    1585 4687.608    lm 1189
## 2  8325.985    3702 4623.985    lm 1208
## 3  8250.050    4267 3983.050    lm 5726
## 4  6466.103    2638 3828.103    lm 5751
## 5  5744.796    1957 3787.796    lm  942

```

```

n_obs <- which(houses$sqm_price == 1585)

```

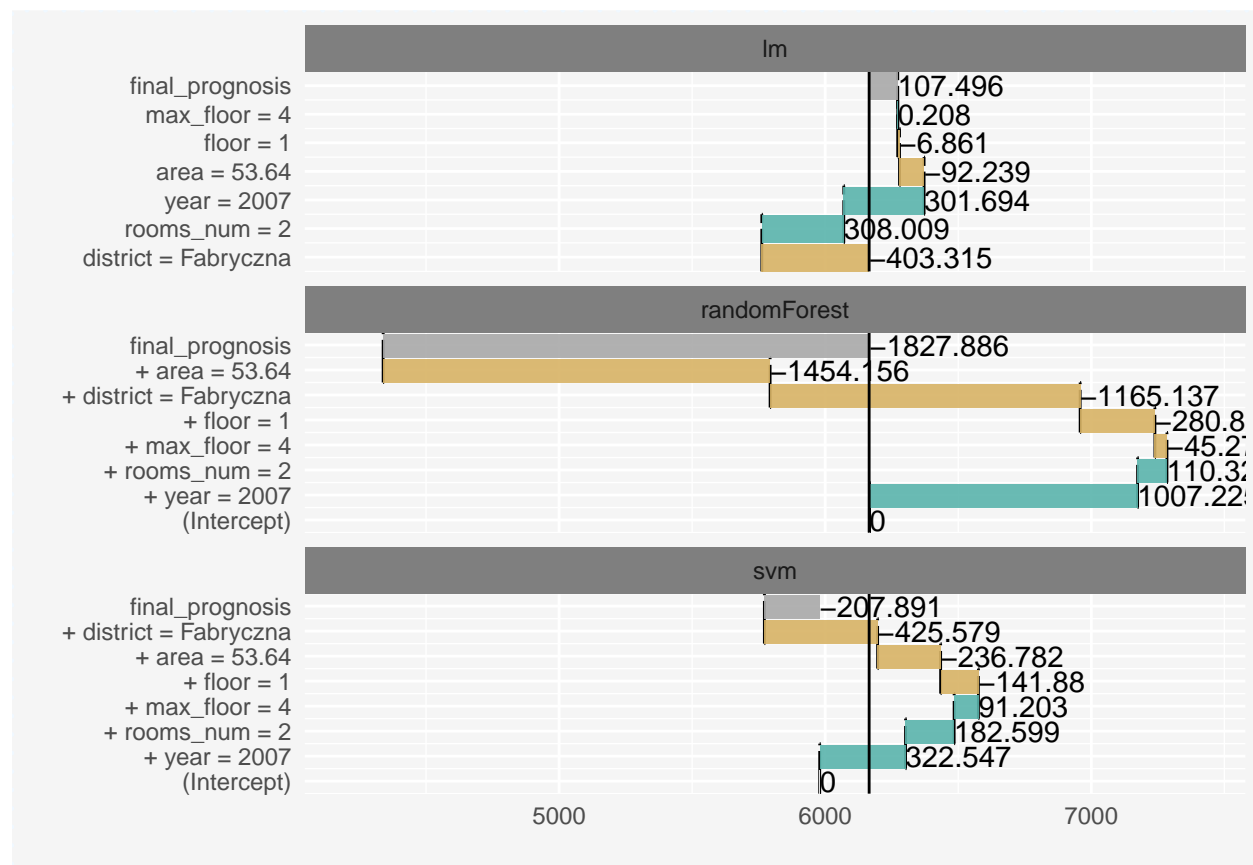
Exercise 2

```

rf_expl_sp <- single_prediction(rf_expl, houses[n_obs, -3])
svm_expl_sp <- single_prediction(svm_expl, houses[n_obs, -3])
lm_expl_sp <- single_prediction(lm_expl, houses[n_obs, -3])

plot(rf_expl_sp,
     svm_expl_sp,
     lm_expl_sp)

```



Exercise 3

```
library(shapleyr)

## Loading required package: checkmate
## Loading required package: combinat
##
## Attaching package: 'combinat'
## The following object is masked from 'package:utils':
##
##      combn
## Loading required package: shiny
## Loading required package: shinydashboard
##
## Attaching package: 'shinydashboard'
## The following object is masked from 'package:graphics':
##
##      box
## Loading required package: reshape2
##
## Attaching package: 'reshape2'
## The following object is masked from 'package:tidyr':
##
##      smiths
## Welcome to the ShapleyR package!
shapley_vals <- shapley(n_obs, house_task, house_rf_mlr)

gather(shapley_vals$values, "colname", "contribution") %>%
  filter(colname %in% colnames(houses)) %>%
  mutate(contribution = as.numeric(contribution)) %>%
  arrange(desc(abs(contribution)))


##      colname contribution
## 1 district      -535.795
## 2 rooms_num    -292.509
## 3 max_floor    -250.853
## 4      year      211.827
## 5      area    -206.529
## 6      floor    -85.600

# alternatively use just
shapley_vals

## $task.type
## [1] "regr"
##
## $feature.names
## [1] "rooms_num" "area"      "year"      "floor"      "max_floor" "district"
```

```
## $predict.type
## [1] "response"
##
## $prediction.response
## [1] 4310.482
##
## $data.mean
## [1] 6165.112
##
## $values
##   _Id _Class rooms_num    area    year floor max_floor district
## 1 1189    NA  -292.509 -206.529 211.827 -85.6  -250.853 -535.795

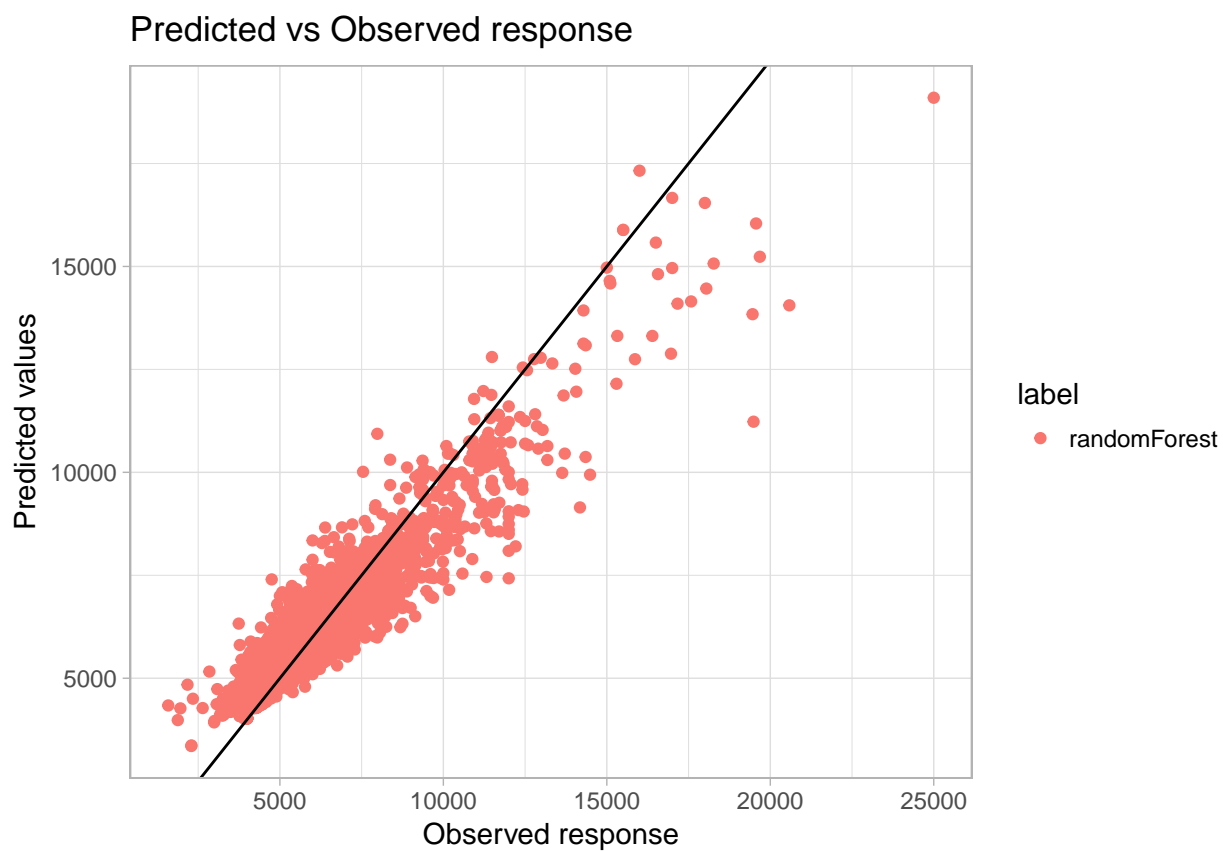
class(shapley_vals) <- c("shapley.singleValue", "list")
plot(shapley_vals)
```



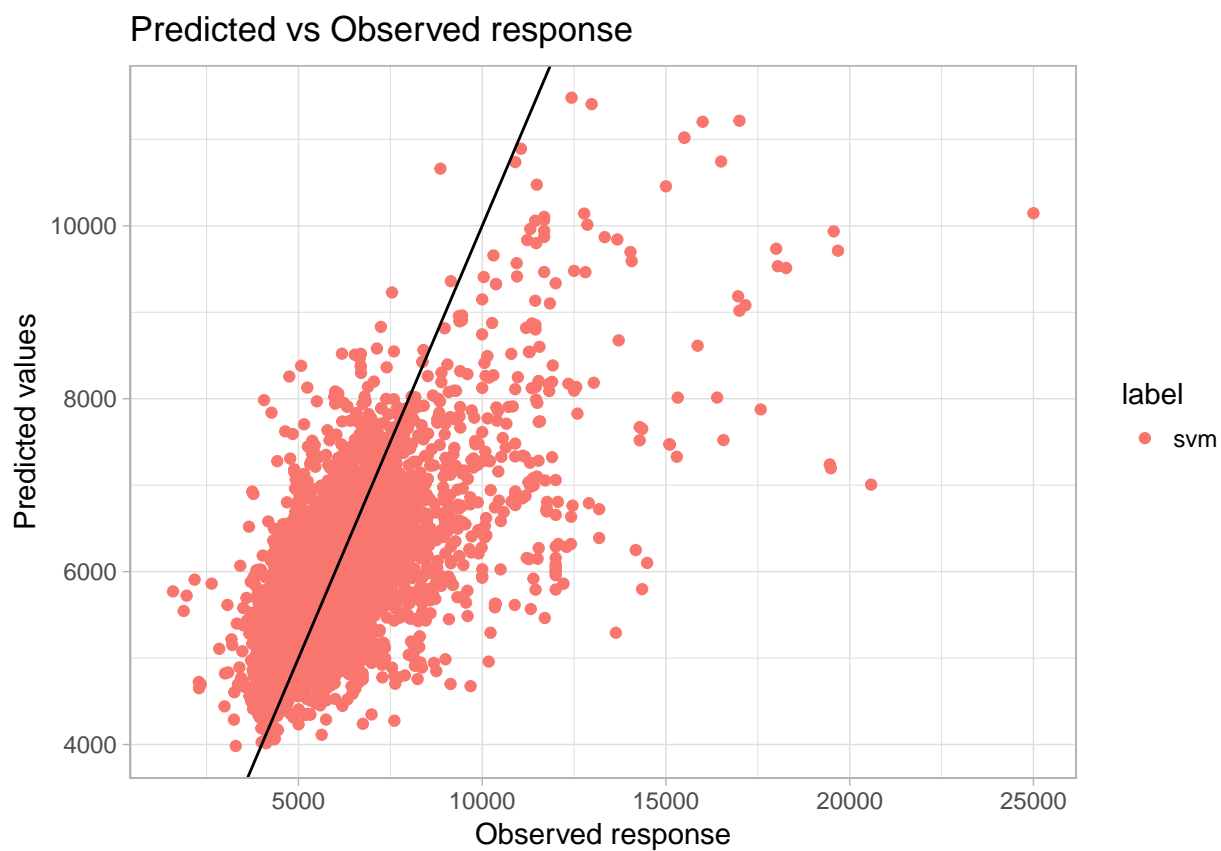
Bonus 1

```
rf_audit <- audit(rf_expl)
svm_audit <- audit(svm_expl)
lm_audit <- audit(lm_expl)

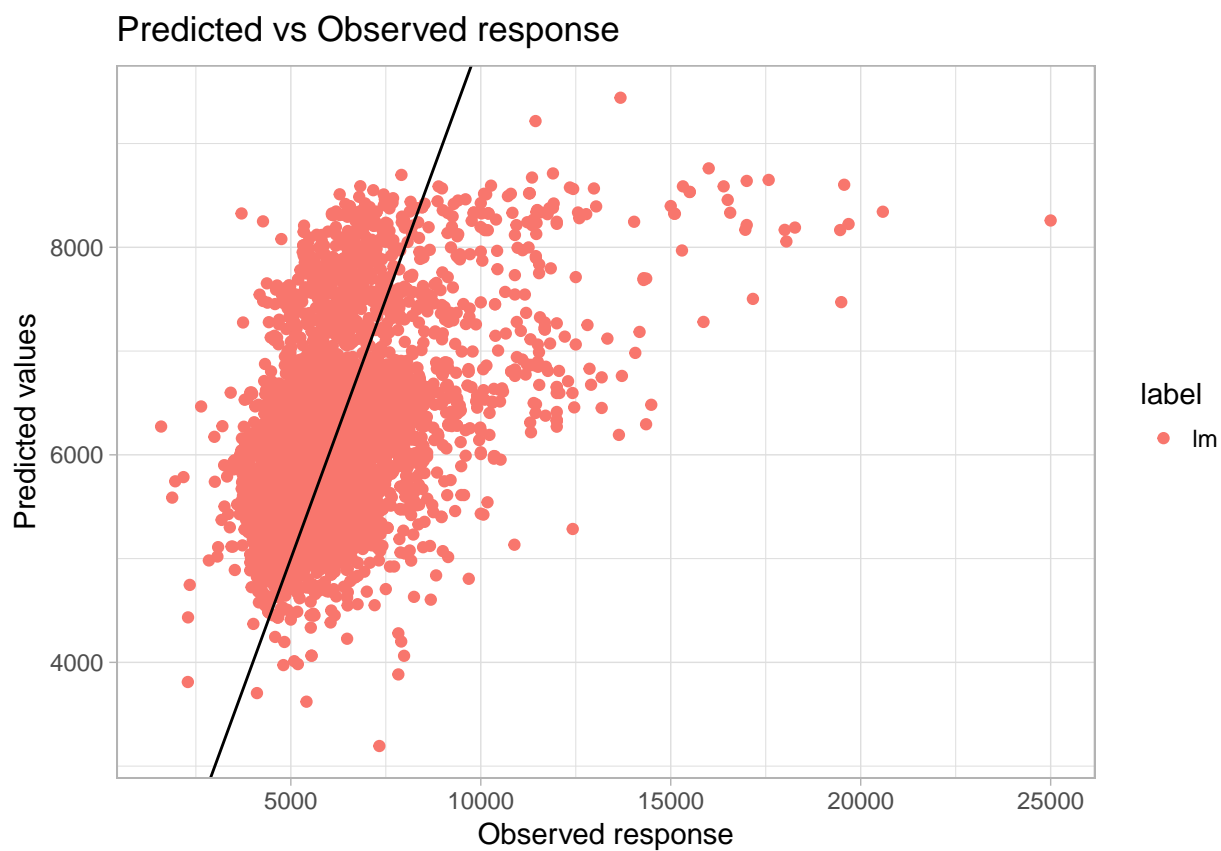
plotPrediction(rf_audit)
```



```
plotPrediction(svm_audit)
```

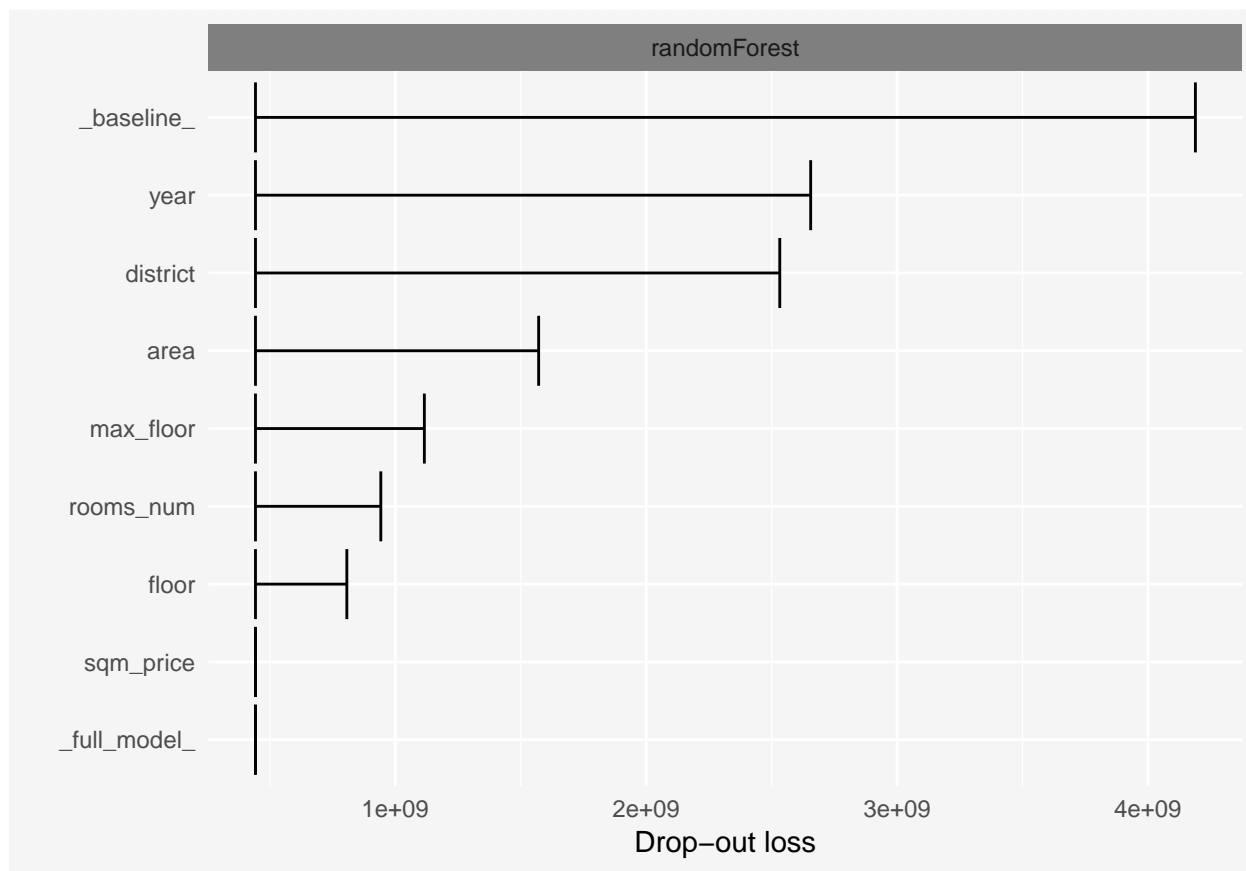



```
plotPrediction(lm_audit)
```



Bonus 2

```
rf_global <- variable_importance(rf_expl)
plot(rf_global)
```



Exercise 4

```
library(live)
library(mlr)

houses_similar <- sample_locally2(houses, houses[n_obs, ], "sqm_price", 1000)
houses_similar2 <- add_predictions2(houses_similar, house_rf)
lm_approx <- fit_explanation2(houses_similar2, "regr.lm")

lm_approx

## Warning in summary.lm(model_tmp): essentially perfect fit: summary may be
## unreliable

## Dataset:
## Observations: 1000
## Variables: 7
## Response variable: sqm_price
## Explanation model:
## Name: regr.lm
## Variable selection wasn't performed
## Weights present in the explanation model
## R-squared: 0.8633
```

Exercise 5

```
plot_explanation2(lm_approx, regr_plot_type = "forest")
```

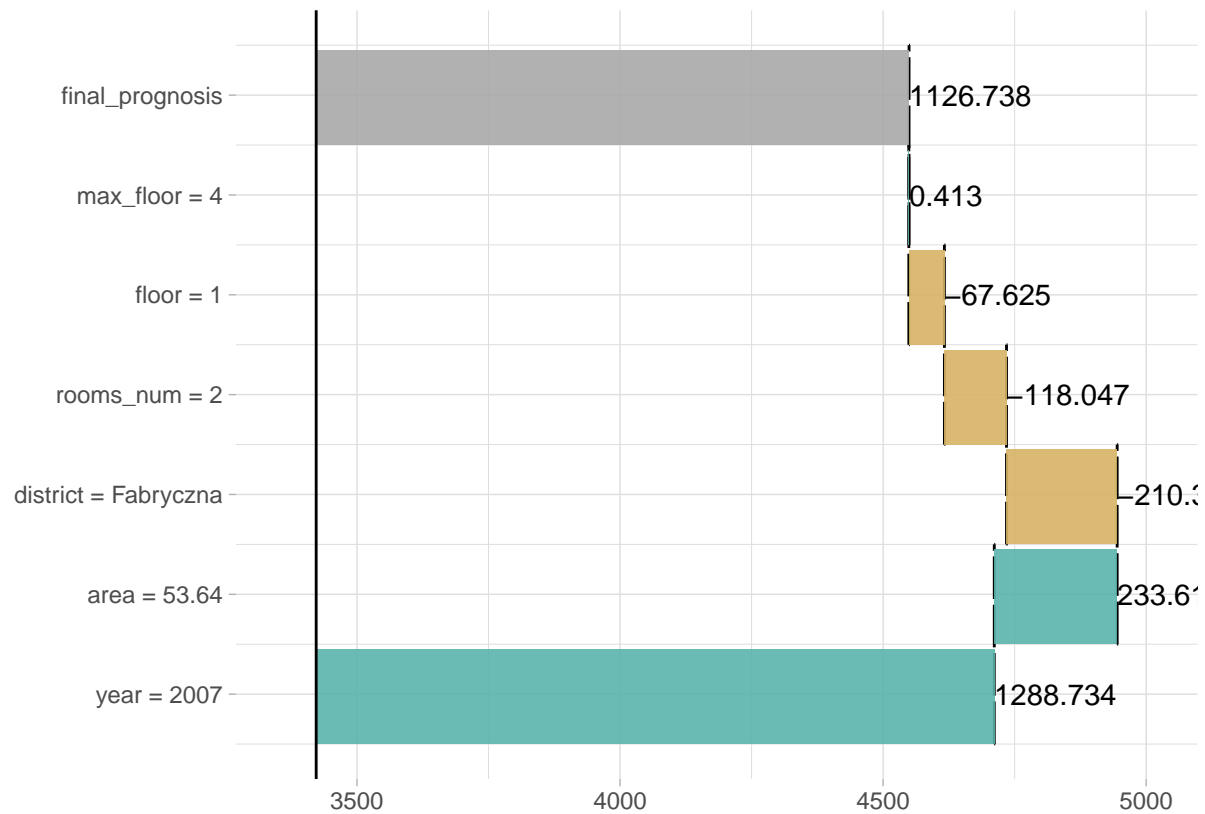
```
## Warning in summary.lm(x): essentially perfect fit: summary may be
## unreliable

## Warning in summary.lm(object): essentially perfect fit: summary may be
## unreliable

## Warning in recalculate_width_panels(panel_positions, mapped_text =
## mapped_text, : Unable to resize forest panel to be smaller than its
## heading; consider a smaller text size
```

Variable	N	Estimate	p
rooms_num	1000	1008.95 (935.05, 1082.85)	<0.001
area	1000	-117.49 (-308.33, 73.35)	0.2
year	1000	263.11 (115.35, 410.88)	<0.001
floor	1000	316.01 (236.02, 395.99)	<0.001
max_floor	1000	4.75 (-88.67, 98.18)	0.9
district	Fabryczna 895	Reference	
	Krzyki 48	1351.43 (1255.57, 1447.28)	<0.001
	Psie Pole 18	973.57 (820.42, 1126.71)	<0.001
	Srodmiescie 26	2309.32 (2181.13, 2437.50)	<0.001
	Stare Miasto 13	5224.04 (5044.50, 5403.57)	<0.001
(Intercept)		-519572.20 (-816320.12, -220824.12)	<0.001

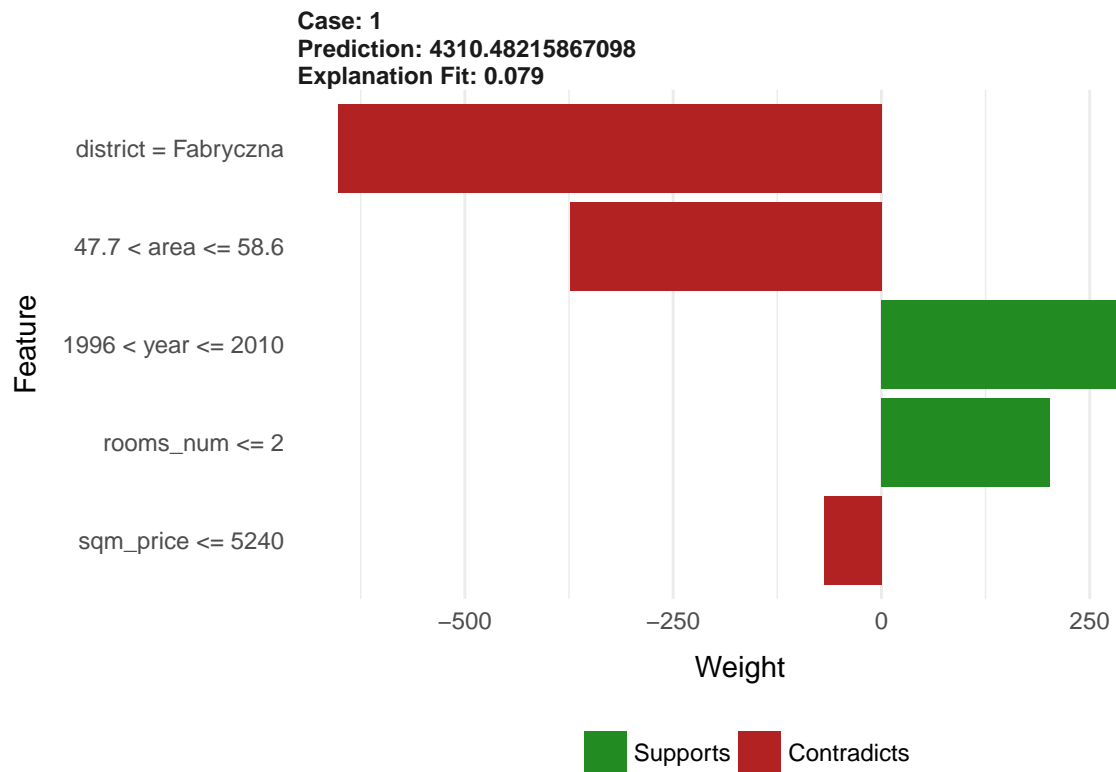
```
plot_explanation2(lm_approx, regr_plot_type = "waterfall")
```



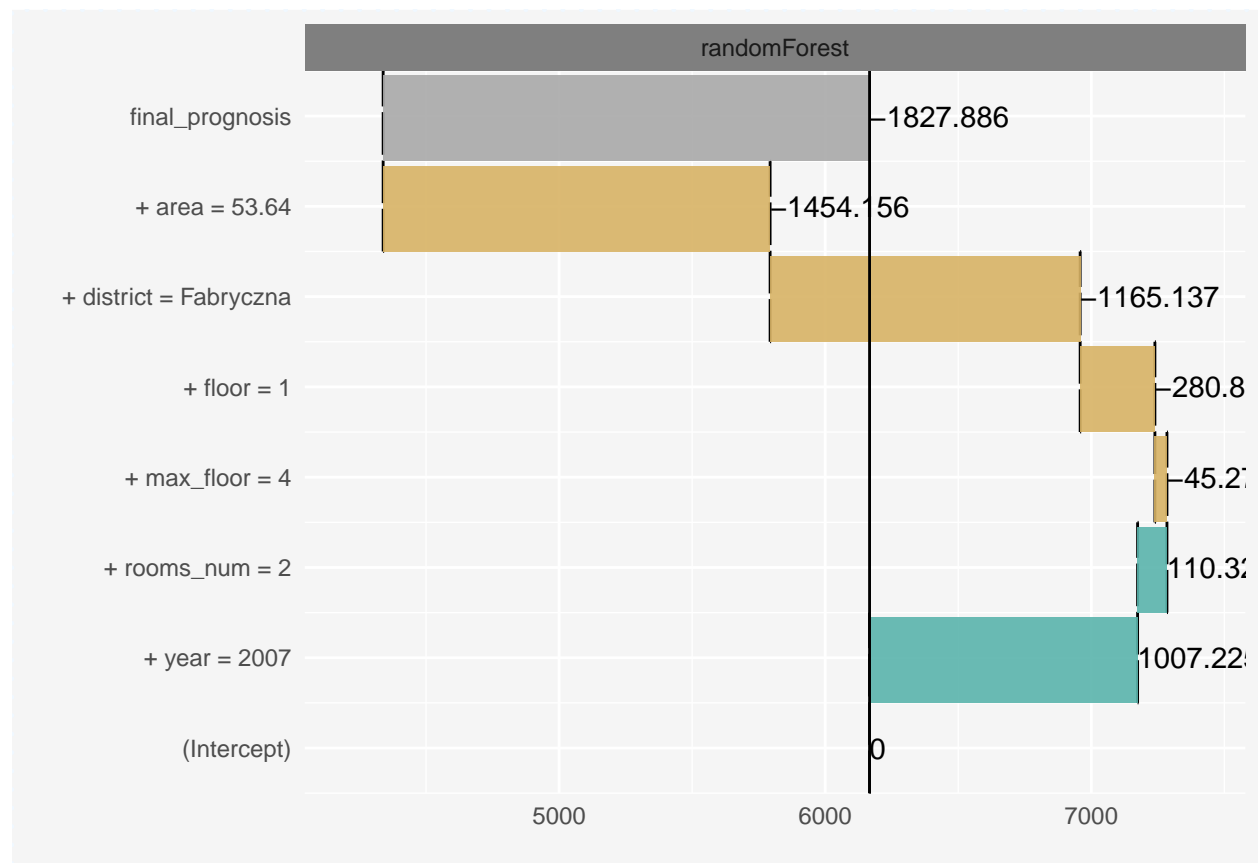
Exercise 6

```
library(lime)

##
## Attaching package: 'lime'
## The following object is masked from 'package:DALEX':
##
##   explain
## The following object is masked from 'package:dplyr':
##
##   explain
lime_rf <- lime(houses, house_rf_mlr)
lime_explanation <- lime::explain(houses[n_obs, ], lime_rf, n_features = 5)
plot_features(lime_explanation)
```



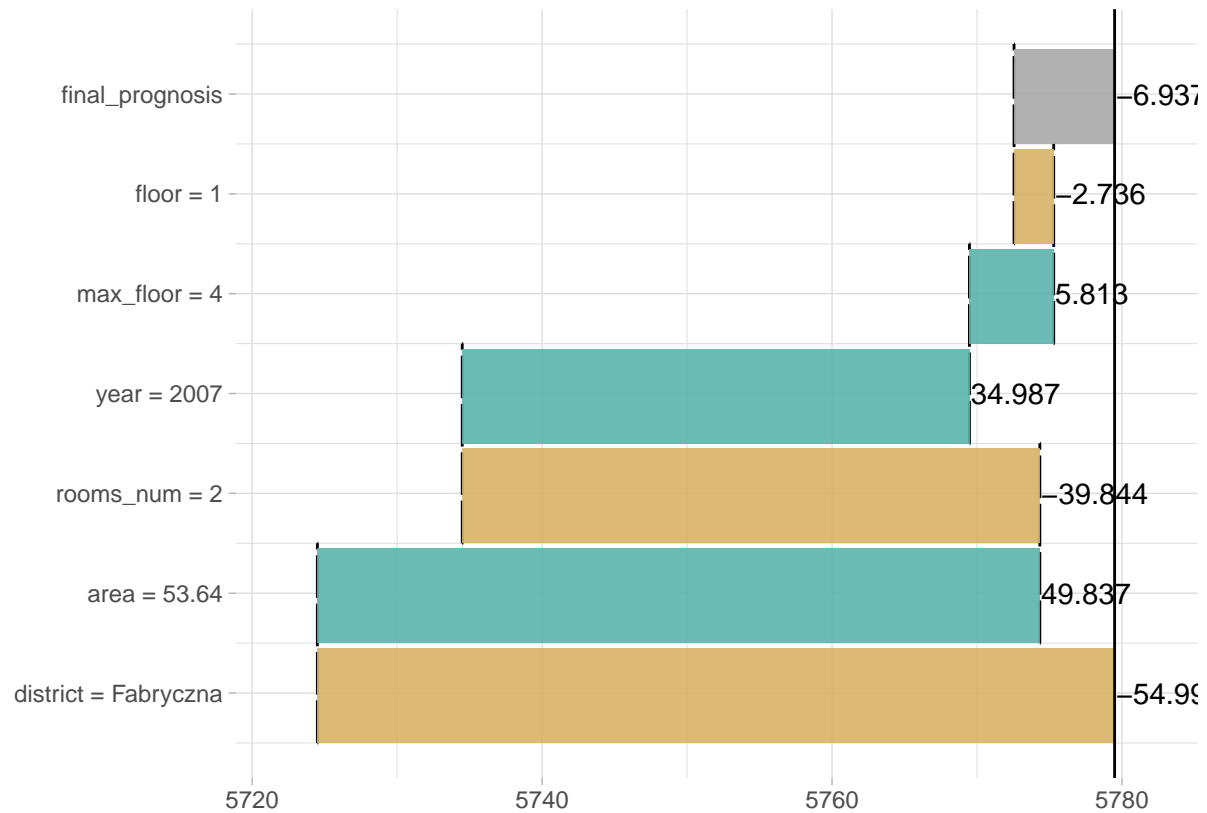
```
plot(rf_expl_sp)
```



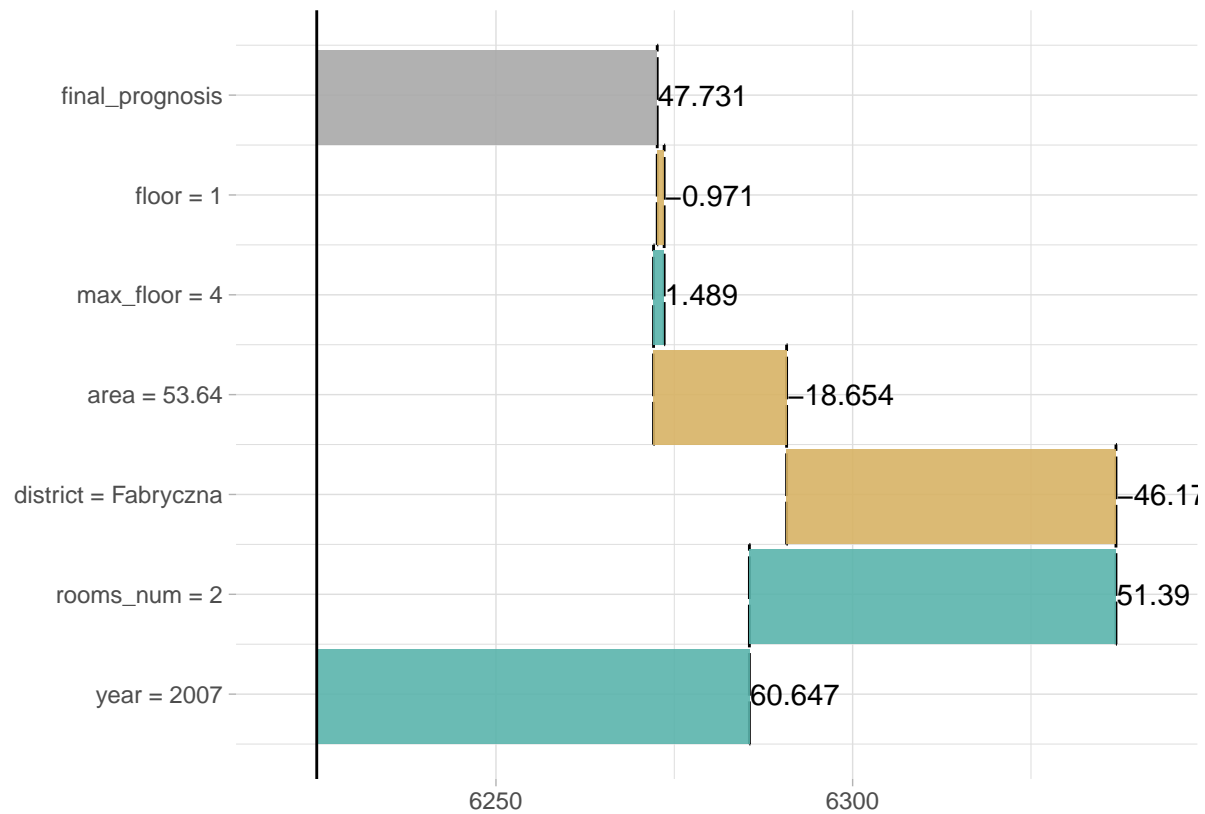
Bonus 3

```
houses_similar3 <- add_predictions2(houses_similar, house_svm)
houses_similar4 <- add_predictions2(houses_similar, house_lm)
lm_approx2 <- fit_explanation2(houses_similar3)
lm_approx3 <- fit_explanation2(houses_similar4)

plot_explanation2(lm_approx2, "waterfall")
```



```
plot_explanation2(lm_approx3, "waterfall")
```



```
plot_explanation2(lm_approx2, "forest")
```

```
## Warning in summary.lm(x): essentially perfect fit: summary may be
## unreliable

## Warning in summary.lm(object): essentially perfect fit: summary may be
## unreliable

## Warning in recalculate_width_panels(panel_positions, mapped_text =
## mapped_text, : Unable to resize forest panel to be smaller than its
## heading; consider a smaller text size
```


Variable	N	Estimate	p
rooms_num	1000	■ -439.23 (-439.23, -439.23)	<0.001
area	1000	■ 9.38 (9.38, 9.38)	<0.001
year	1000	■ 12.38 (12.38, 12.38)	<0.001
floor	1000	■ 4.54 (4.54, 4.54)	<0.001
max_floor	1000	■ 17.12 (17.12, 17.12)	<0.001
district	Fabryczna 895	■ Reference	
	Krzyki 48	■ 277.96 (277.96, 277.96)	<0.001
	Psie Pole 18	■ -412.57 (-412.57, -412.57)	<0.001
	Srodmiescie 26	■ 569.83 (569.83, 569.83)	<0.001
	Stare Miasto 13	■ 1956.82 (1956.82, 1956.82)	<0.001
(Intercept)		■ -18275.81 (-18275.81, -18275.81)	<0.001

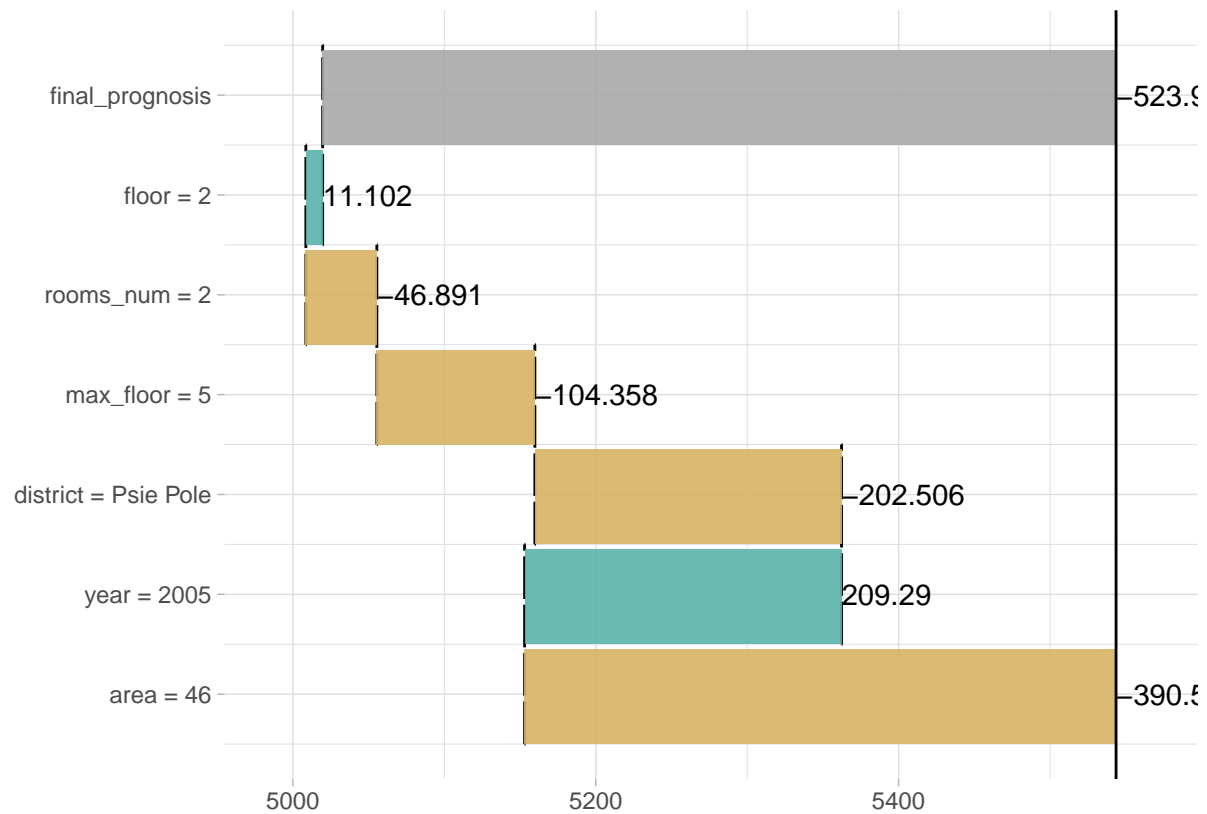
Bonus 4

```
library(live)
library(mlr)
n_obs2 <- 5830
houses_similar <- sample_locally2(houses, houses[n_obs2, ], "sqm_price", 1000)
houses_similar2 <- add_predictions2(houses_similar, house_rf)
n_obs2expl <- fit_explanation2(houses_similar2)
n_obs2expl

## Warning in summary.lm(model_tmp): essentially perfect fit: summary may be
## unreliable

## Dataset:
## Observations: 1000
## Variables: 7
## Response variable: sqm_price
## Explanation model:
## Name: regr.lm
## Variable selection wasn't performed
## Weights present in the explanation model
## R-squared: 0.8071

plot_explanation2(n_obs2expl, "waterfall")
```



```
plot_explanation2(lm_approx, regr_plot_type = "waterfall")
```

